

A Probabilistic Approach to Content-Based Audio Analysis for Music Clustering and Search

NADER AL-NAJI

LAWRENCE DIAO

JOSH CHEN

Department of Computer Science
Princeton University

Abstract

x

Pupko02 method:

$P[X|r] = \text{sum}(\text{internal node assignments})$ distance from internal node to external node Max likelihood, so find r that maximizes this.

Mayrose04 EB method:

Let $P[X|r]$ be as in Pupko02

Approximate $P[r]$ as gamma distribution with 16 bins, parameters $\alpha, \beta=1$? Each bin has value equal to $1/2(\text{gamma}(\text{left}) + \text{gamma}(\text{right}))$, and has boundaries such that $1/(\text{num bins})$ of the PDF falls into the bin.

$P[X]$ is just marginalized over all bins

Find $E[r|X]$. MAP doesn't seem to make sense.

Analysis using human-labeled textual metadata for songs has become a technique widely used in many industrial-strength music analysis systems. Such methods are appealing because they allow one to make full use of a large corpus of text-based information retrieval literature. However, textual labels can be very subjective, depending on who is doing the labeling, and such schemes inherently leave behind a lot of useful information by ignoring the signal itself. By analyzing the actual audio signal, so-called content-based approaches have the potential to pick up on what text-based approaches miss. This paper will discuss and build upon previous research by Perry Cook et al. [3] in the relatively new field of content-based audio analysis.

Our analysis is centered around the idea that songs can be represented as probability distributions over feature space. Using this idea, we

then seek a similarity measure $\text{dist}(A, B)$ (not to be confused with a similarity *metric*) between two songs A and B , modeled as probability distributions, such that $\text{dist}(A, B) < \text{dist}(A, C)$ if and only if song A is perceptually more similar to B than to C . Of course, because the performance of such a similarity measure will rely heavily on the methods used to extract features, a large part of the paper will be devoted to outlining exactly how features are extracted from songs.

Establishing a similarity measure between songs allows us to evaluate our model using standard information retrieval metrics. Taking one song out of a collection of songs, we then "search" for the songs closest to that song based on our similarity measure. We then compare the results of this search to what songs would be most similar perceptually.

Finally, we present a novel method we developed to cluster music under this probabilistic framework, and show the results of clustering using this method on an actual collection of songs.

I. FEATURE EXTRACTION

Computing content-based similarity between songs relies heavily on extracting relevant features from audio. Thus, we detail the exact steps taken to condense a piece of music into a large collection of feature vectors, to be used in further analysis. Mel-Frequency Cepstral

Coefficients, or MFCC's, have been shown by Logan [1] to be a reasonable way to extract perceptually relevant features from music as well as speech. The methods described here for completeness are thus the standard implementation of MFCC's described in [1] applied to a whole song.

Song Feature Extraction Steps

1. **Break song up into 50ms windows with a hop size of 20ms.**

The reasons for using these time frames are described below.

2. **For each window:**

- **Apply a Hamming window to samples**
This amounts to smoothing the window with a "raised cosine" function to reduce edge effects when taking the FFT.

- **Take the Fast Fourier Transform of the window**

This turns the song window into a vector of complex values, each corresponding to a frequency. If the song was sampled at 44.1kHz (a common sampling rate) and a window size of 50ms is used, then the maximum frequency that will have a coefficient in the resulting vector will be 22kHz, just at the limit of what humans can hear. The windows overlap to compensate for the fact that we apply a Hamming window to eliminate edge effects.

- **Drop the phase and log the amplitudes**
We drop the phase because, as mentioned in [1], phase values have been shown to be perceptually insignificant. We log the amplitudes because humans perceive volume logarithmically, hence the Decibel scale.

- **Bin the frequencies according to the Mel Scale**

Humans perceive increases in pitch non-linearly with increases in frequency. For example, an increase in one octave from middle C to high C corresponds to a *doubling* of frequency. The Mel Scale maps frequencies to Mel values such that linear increases in Mel values represent linear increases in perceived pitch. Thus, at this step we take a set of mel-spaced bins and populate them with the average over their corresponding frequency ranges.

- **Take the Discrete Cosine Transform of the resulting Mel spectrum**

This final step is done as a form of principal component analysis to compress the Mel vector into a lower-dimensional vector with minimal loss of information. This works because the DCT concentrates most of the information of a given signal into a few low-frequency components, approaching the KL Transform which is optimal in terms of loss of information [1].

- **Take the first 20 coefficients of the DCT to be a 20-dimensional feature vector**

The other DCT coefficients can be dropped with minimal loss of information since, as mentioned, most of the Mel vector's information will be concentrated into the lower-frequency components after applying the DCT. We discard the 0th component of the DCT because it corresponds to volume, to which we want to be invariant.

For each song we end up with a large number of 20-dimensional feature vectors, where each vector corresponds to a particular window of time. Figure I illustrates feature extraction for a single window.

II. SONGS AS PROBABILITY DISTRIBUTIONS

After feature extraction, every song has a large number of 20-dimensional feature vectors (on the order of thousands) corresponding to it, with each vector representing a specific window within the song. Considering songs simply as a list of ordered feature vectors makes them difficult to deal with as well as difficult to store. In fact, this representation under a naive implementation could require more space to store than the actual song. Thus, we make some careful simplifying assumptions in order to be able to condense this large number of feature vectors into a more meaningful representation that is also easier to deal with.

The first simplifying assumption we make is that feature vectors within a song can be reordered with only a small loss in meaningful information. Recall that one of our main goals is to develop a similarity measure with which to compare songs, so whether or not songs remain similar upon reordering of their windows is

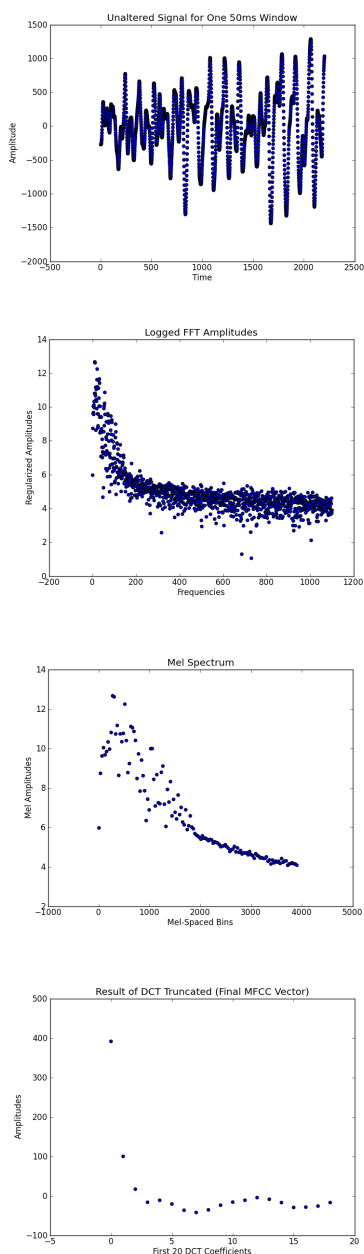


Figure I: An illustration of the feature extraction process.

crucial. This assumption certainly causes our model to neglect some useful information, namely how a song evolves over time. However, it is not unreasonable to think that two similar

songs, which presumably contain many feature vectors with features unique to their genre in common, could still be deemed similar even if order is disregarded. For example, one can still tell that two songs by Metallica are similar, even if they only hear random chunks out of order, since random chunks from both of these songs will contain features unique to heavy metal music with high probability. This "bag of windows" assumption thus appears to be justified in the sense that similarity between songs, while not completely unscathed, is preserved to a sufficient extent.

Next, assuming songs can be thought of as bags of windows, we consider a generative probability model for songs, which allows us to represent each song using only the parameters for its probability distribution. This is useful since this parameterized distribution, by virtue of being a mathematical object, is much easier to deal with (and store) than thousands of unordered feature vectors. In particular, we assume that these unordered vectors are generated by an underlying Mixtures of Gaussians distribution.

The reasoning behind the Mixtures of Gaussians assumption is as follows. A song typically consists of multiple sections: an intro, a chorus, etc... and each of these sections will have feature vectors corresponding to it. Further, vectors from a particular section will likely contain features unique to that section. For example, vectors corresponding to the intro of a Metallica song might all contain features unique to guitar solos, while vectors from the chorus might all contain features unique to the lead singer's voice. We thus consider songs as generated by an underlying Mixtures of Gaussians distribution where each mixture corresponds to a kind of part of a song. Songs themselves, then, under the bag of windows assumption, can be seen as generated by choosing a mixture from a particular song's distribution according to its prior, sampling a vector from that mixture, and repeating this process. Figure II illustrates the generative model:




Figure II: The generative model for a song.

Having justified both the bag of windows assumption and the generative model, we compute the latent parameters for each song's distribution using standard Mixtures of Gaussians Expectation Maximization. Computing the cluster priors, means, and covariances, then, gives us a full parameterization of each song and allows us to consider songs as probability distributions over feature space, instead of simply as bags of unordered feature vectors. The former characterization is much easier to deal with and opens up a range of possibilities, affording us the flexibility of comparing songs by comparing their probability distributions.

III. A PROBABILISTIC SIMILARITY MEASURE

We compute a Mixtures of Gaussians probability distribution for each song using EM as discussed above. Now, having extracted a parameterized probability distribution over feature space for each song, we discuss several ways to compare these distributions and, by proxy, compare the songs.

We first naturally gravitated towards a similarity measure very much inspired by Euclidean distance. We define $dist(A, B)$ as the square of the difference between the two distributions integrated over every point in feature space. To be clear, taking f and g to be probability distributions over n -dimensional feature space corresponding to songs A and B respectively, we would want to compute $dist(A, B)$ as:

$$\int_{-\infty}^{\infty} \dots \int_{-\infty}^{\infty} (f(x_1, \dots, x_n) - g(x_1, \dots, x_n))^2 dx_1 \dots dx_n. \quad (1)$$

In theory this provides us with the most intuitive measure of similarity or distance between arbitrary distributions. This similarity measure also has the added benefit of being a metric since $dist(A, A) = 0$, $dist(A, B) = dist(B, A)$, and it obeys the triangle inequality. However, in practice, even with only twenty-dimensional features, $n = 20$, and using MoG's as our distributions for f and g , this sum is difficult to compute using numerical methods and doesn't appear to have a closed form (though it could).

Since this distance metric seems to be difficult to compute when dealing with Mixtures of Gaussians, we resort to approximation. We consider the following unbiased estimator of (1) that consists of sampling points from each distribution and computing the normalized squared difference between the two distributions at each point. In the following equation, f and g are again probability distributions over n -dimensional feature space, and x_i and y_i are n -dimensional points sampled from f and g respectively, with N being the total number of samples:

$$\frac{1}{N} \left(\sum_{x_i} \frac{(f(x_i) - g(x_i))^2}{f(x_i)} + \sum_{y_i} \frac{(f(y_i) - g(y_i))^2}{g(y_i)} \right) \quad (2)$$

Given a large enough number of samples, this would theoretically approach the metric described above. However, in practice this metric is not only very slow to compute under our implementation (using $N = 2000$), but also the results of using this similarity metric to compare songs are sub-par compared to the results of using the method we describe next. It is possible that even though this estimator is unbiased, its variance is too large to make using only 2000 points in the sample reasonable. In future research we want to explore this metric in further detail using a larger number of samples, after speeding up our implementation.

The similarity approximation actually used was a simple Monte Carlo approximation to Kullback-Leibler divergence, a non-symmetric measure of the difference between two probability distributions widely used in many aspects of speech and image recognition. We approximate the exact Kullback-Leibler measure for

continuous densities f and g :

$$D(f||g) = \int_{-\infty}^{\infty} f(x) \ln \frac{f(x)}{g(x)} dx \quad (3)$$

using Monte Carlo sampling as suggested in [2]. Taking f and g to again be the densities of two MoG distributions corresponding to songs, we sample N iid points, $\{x_i\}_{i=1}^N$, from f and compute:

$$D_{MC}(f||g) = \frac{1}{N} \sum_{i=1}^N \log \frac{f(x_i)}{g(x_i)}. \quad (4)$$

This additionally turns out to be an unbiased estimator of KL divergence since:

$$E_f\left[\frac{1}{N} \sum_{i=1}^N \log \frac{f(x_i)}{g(x_i)}\right] = \int_{-\infty}^{\infty} f(x) \ln \frac{f(x)}{g(x)} dx = D(f||g) \quad (5)$$

Further, since we are simply taking the average over iid samples, the Central Limit Theorem applies and we see that $Var[D_{MC}] = \frac{1}{N} Var_f[\log f/g]$, meaning the variance of our estimator decreases as N increases, as discussed in [2]. This measure also captures the intuition that songs should be more likely to generate points that came from songs similar to themselves.

We use this approximation without the $f(x_i)$ term when doing music retrieval, since f doesn't change when searching for a specific song, to give us the following expression:

$$-\frac{1}{N} \sum_{i=1}^N \log g(x_i). \quad (6)$$

We do the same when using KL divergence to cluster songs.

IV. PROBABILISTIC MUSIC RETRIEVAL

For this section, we define $dist(A, B) = -\frac{1}{N} \sum_{i=1}^N \log g(x_i)$, where g is the probability distribution corresponding to song B , and the x_i are points sampled from f , the distribution corresponding to song A . We test this measure by doing music retrieval searches, returning a list of the songs in our collection ordered by their distance from the query song, using a library of

approximately 160 songs. We then evaluate the relevance of the results using standard information retrieval metrics. We acknowledge that the library of songs we use is relatively small, containing songs from only about thirty artists, but we emphasize the accuracy of the results even on this small dataset, as well as the scalability of our retrieval computation (It is "embarrassingly parallelizable.").

Figure III at the end of this paper illustrates the results of queries for selected songs using our measure and our 160-song library. The results returned when querying for a song A are ordered by their similarity to A based on the measure discussed.

Note how the top ten songs returned for a given query always consist mostly of songs with the same artist as the query song, and that, overall, songs returned that are not by the same artist are songs that are perceptually similar to the query. For example, the top ten results in a query for a Steely Dan song are all Steely Dan, and the eleventh result is a song by Eric Clapton, another prolific rock artist. This pattern is prevalent in all of the searches performed, even across different genres, showing that perceptually significant features from all genres are being captured by our system. Further, songs from similar genres and bands are confused in a way consistent with how humans would confuse them. For example, in a query for "I Want It That Way" by the Backstreet Boys, "Tearin' Up My Heart", a song by 'N Sync, a very similar boyband, is returned above "The Perfect Fan", another song by the Backstreet Boys.

We also include an evaluation of how this system performs using standard information retrieval metrics, shown in the figure above. We compute the reciprocal rank, precision at rank 15, discounted cumulative gain at Rank 15, and the average precision at each point a song in the same genre appears for each of the queries above. As can clearly be seen, this system's precision is very high with almost 90% of songs returned in the same genre. Further, the classical music query reaches the maximum DCG value of 194.7, with the average DCG not too far below that.

V. A TECHNIQUE FOR CLUSTERING SONGS

Clustering data provides insight into what features of the data are being captured by a particular model. Hence, we now try clustering our data using the full (symmetrized) KL divergence approximation as our distance measure:


$$\text{dist}(A, B) = \frac{1}{N} \sum_{i=1}^N \log \frac{f(x_i)}{g(x_i)} + \frac{1}{N} \sum_{i=1}^N \log \frac{g(x_i)}{f(x_i)} \quad (7)$$

where f and g are the MoG distributions for songs A and B respectively. Note that we symmetrize the measure so that $\text{dist}(A, B) = \text{dist}(B, A)$.

Since we have a clearly defined similarity measure, the first thing we try is hierarchical clustering. We perform hierarchical agglomerative clustering, using the median similarity criterion, on the songs using a similarity matrix with each value computed according to the approximation above. Figure IV is the dendrogram obtained. Looking at this dendrogram, we see noticeable decreases in height for about ten to fifteen splits. Thus, we generate a flat clustering using $K = 15$ based on this hierarchy. Some of the more significant clusters are shown in figure V at the end of this paper.

From the clusters shown in figure V, we can clearly see that hierarchical clustering is creating groups based on perceptual similarities. It finds and groups all of the classical music

Figure IV: Dendrogram found using median criterion.



in the collection, as well as all of the R&B. Classic rock and pop (namely boy band songs) are grouped as well. What isn't shown is a small set of singleton clusters and one large absorbing cluster, containing a hodge-podge of songs from different genres, that was too large to include. In an effort to improve these flaws, we propose an alternative clustering method based on Expectation Maximization.

The K-Means algorithm is a widely accepted way to cluster related pieces of data. However, a direct application of the K-Means algorithm doesn't quite make sense in our framework since it is not clear what the "mean" of a collection of MoG distributions is. However, if we take the bag of windows assumption seriously, that is, assuming songs can really be seen as unordered bags of windows represented by feature vectors, then we can come up with a natural Expectation Maximization algorithm that emulates K-Means.

The fundamental idea is this: since clusters contain songs, and songs are bags of windows, we can consider a cluster to be a bag of windows as well (essentially a bag of bags of windows). A collection of windows in a cluster, then, can be seen as generated by choosing one of its MoG components uniformly at random, sampling from that component, and then repeating this process. So, we define a cluster's distribution to be the uniform mixture of its MoG components. Further, by considering a cluster to be a probability distribution over feature space, we can then approximate the similarity between a song and a cluster using KL divergence.

We now describe an EM algorithm to cluster

songs where K is the number of initial clusters, I is the number of iterations, and N is the number of points to sample for each song:

Clustering Algorithm(K, I, N):
Presample N points for each song
Precompute $\text{dist}(A, B)$ for all songs
Assign each song to a cluster randomly
for $i \leftarrow 1$ to I
 for each song A
 for each cluster k
 (1) **Compute $\text{dist}(A, k)$**
 (2) **Assign song A to closest cluster**
return clustering

In (1) above, we compute the asymmetric approximation to KL Divergence. That is, we compute:

$$\text{dist}(A, k) = -\frac{1}{N} \sum_{i=1}^N \log g(x_i) \quad (8)$$

where g is the distribution for cluster k and the x_i are sampled from song A 's distribution. This is essentially the probability that cluster k generated song A , and it can be computed very quickly using the log-sum-exp trick, to sum log likelihoods over the MoG components of each cluster, and the precomputed values for song similarity. Here we again leave out the $f(x_i)$ term since f doesn't change when comparing different clusters to the same song. In (2) we assign each song to the cluster with the lowest distance value; that is, we assign each song to the cluster that was most likely to generate it.

Since each cluster is essentially a uniform mixture of the songs contained within it, it appears reasonable to consider this mixture to be the "mean" distribution for the cluster of songs. Considering the cluster distribution in this way, it should be clear that this approach is nothing more than a generalization of the K-Means algorithm to handle probability distributions as input data.

Figure VI at the end of this paper shows a subset of the results of running this algorithm on our collection of songs with $K = 10$. It is worth noting that the clusters generated by this algorithm are significantly better than the clusters generated using hierarchical clustering. Here, not only are there no singleton clusters, but every cluster clearly represents a genre and, a lot of the time, a specific artist. While this method

is more impacted by initialization (a different clustering is obtained depending on the initial clusters used to seed the algorithm), by starting each cluster off with a single random song from the collection, we found the results almost always converge to the same clustering.

VI. FUTURE WORK

While we covered a lot of ground in this paper, we still feel as though there is much work to be done. For one, the feature extraction mechanism we implemented was only the bare minimum of what we actually wanted to do. While the MFCC feature extraction technique appeared to work very well, in the future we want to adapt MFCC's, which were developed specifically for speech recognition, to make them better at representing musical data. For example, possibly because speech recognition doesn't rely heavily on pitch, MFCC's appear not to take changes in perceived pitch into account, even though humans perceive some pitches much more strongly than others. We also wanted to experiment with MFCC flux, that is the difference between MFCC vectors at each point, to see how adding first-order derivative vectors affects our results. We wanted to experiment with using distributions other than Mixtures of Gaussians to cluster the MFCC vectors and see how this affects the similarity measure. We used mixtures of five Gaussians to cluster our vectors; however, it is not terribly clear that this number is optimal and being able to figure out the best number of mixtures each song should be represented with would also be an important goal. KL Divergence worked very well as a similarity measure; however, the Monte Carlo approach to estimating it forces one to make a trade-off between computation time (number of samples) and accuracy. In the future we want to experiment with more efficient and more accurate approximations to KL Divergence, as well as try completely new measures (and metrics) to see how they perform. Finally, while the bag of windows assumption simplifies a lot of things, in the future we would like to experiment with temporal models to seriously take into account how songs evolve over time, and include this information in our analysis.

VI. REFERENCES

- [1] B. Logan, *Mel Frequency Cepstral Coefficients for Music Modelling*, Cambridge Research Laboratory.
- [2] J. Hershey and P. Olsen, *Approximating the Kullback Leibler Divergence Between Gaussian Mixture Models*, IBM T.J. Watson Center.
- [3] G. Tzanetakis, *Musical Genre Classification of Audio Signals*, IEEE Transactions on Speech and Audio Processing Vol. 10 No. 5, 2002.
- [4] D. Blei, M. Hoffman, and P. Cook, *Content-Based Music Similarity Computation Using the Hierarchical Dirichlet Process*.
- [5] J. Aucouturier and F. Pachet, *Music Similarity Measures: What's the Use?*.

Figure III: Music Retrieval Using Approximation to KL Divergence (Without Normalization)**a) Classical Music (Rachmaninov)**

Song being searched for: 12_-_Prelude_in_C_major,_Op.32.No.1

Songs closest to this song:

- 1: 12_-_Prelude_in_C_major,_Op.32.No.1: -68591.204978
- 2: 24_-_Prelude_in_D_flat_major,_Op.32.No.13: -73033.827607
- 3: 15_-_Prelude_in_E_minor,_Op.32.No.4: -73506.562205
- 4: 06_-_Prelude_in_G_minor,_Op.23.No.5: -73638.926928
- 5: 14_-_Prelude_in_E_major,_Op.32.No.3: -74374.963012
- 6: 07_-_Prelude_in_E_flat_major,_Op.23.No.6: -74412.354153
- 7: 19_-_Prelude_in_A_minor,_Op.32.No.8: -74686.723507
- 8: 04_-_Prelude_in_D_minor,_Op.23.No.3: -74721.104207
- 9: 20_-_Prelude_in_A_major,_Op.32.No.9: -74985.791112
- 10: 08_-_Prelude_in_C_minor,_Op.23.No.7: -75102.159756
- 11: 01_-_Prelude_in_C_sharp_minor,_Op.3.No.2: -75110.30196
- 12: Vladimir_Ashkenazy_-_Prelude_in_C_minor,_Op.23.No.7: -75367.916979
- 13: 21_-_Prelude_in_B_minor,_Op.32.No.10: -75547.950154
- 14: 23_-_Prelude_in_G_sharp_minor,_Op.32.No.12: -75665.925955
- 15: Prelude2: -76398.211258

c) Classic Rock (Eric Clapton)

Song being searched for: 01-Eric_Clapton-Signe

Songs closest to this song:

- 1: 01-Eric_Clapton-Signe: -75922.412186
- 2: 13-Eric_Clapton-Old_Love: -80844.87701
- 3: 05-Eric_Clapton-Lonely_Stranger: -80456.36072
- 4: 06-Eric_Clapton-Nobody_Knows_You_When_You're_Down_&_Out: -81552.032228
- 5: 02-Eric_Clapton-Before_You_Accuse_Me: -81905.231136
- 6: 07-Eric_Clapton-Layla: -82309.644601
- 7: 03-Eric_Clapton-Hey_Hey: -82578.647622
- 8: 12-Eric_Clapton-Malted_Milk: -83142.276544
- 9: 08-Eric_Clapton-Running_On_Faith: -83332.194951
- 10: 04-Eric_Clapton-Fears_In_Heaven: -83920.588263
- 11: 10-Eric_Clapton-Alberta: -84274.192797
- 12: 04_-_It's_Gotta_Be_You: -85153.731215
- 13: SteelyDan2: -85573.667245
- 14: 11-Eric_Clapton-San_Francisco_Bay_Blues: -85626.624126
- 15: 06_Steely_Dan_-_I_Got_The_News: -85732.5736

e) Boy Bands (Backstreet Boys)

Song being searched for: 02_-_I_Want_It_That_Way

Songs closest to this song:

- 1: 02_-_I_Want_It_That_Way: -76024.091635
- 2: 03_-_Show_Me_The_Meaning_Of_Being_Lonely: -78433.01848
- 3: 07_-_Don't_Wanna_Lose_You_Now: -78754.466319
- 4: 06_-_Don't_Want_You_Back: -78808.847484
- 5: 08_-_The_One: -78965.104408
- 6: 09_-_Back_To_Your_Heart: -79464.563236
- 7: 05_5_0'clock_(feat._Lily_Allen_and_Wiz_Khalifa): -79802.186367
- 8: 01_-_Tearin'_Up_My_Heart: -79937.715927
- 9: 12_-_The_Perfect_Fan: -80108.859516
- 10: 04_-_For_The_Girl_Who_Has_Everything: -80276.630152
- 11: 13_-_Giddy_Up: -80319.443261
- 12: 08_-_I_Want_You_Back: -80398.577942
- 13: 09_-_Everything_I_Own: -80518.221767
- 14: 04_-_It's_Gotta_Be_You: -80549.390305
- 15: 06_-_You_Got_It: -80550.72322

g) Pop (Britney Spears)

Song being searched for: 06_-_If_U_Seek_Amy_[Main_Version]

Songs closest to this song:

- 1: 06_-_If_U_Seek_Amy_[Main_Version]: -75929.337897
- 2: 02_-_Circus_[Main_Version]: -79277.894309
- 3: 11_-_Lace_And_Leather_[Main_Version]: -79455.920435
- 4: 13_-_Radar_[Main_Version]: -79711.211247
- 5: 05_-_Shattered_Glass_[Main_Version]: -79711.83541
- 6: 19_-_My_Dad's_Gone_Crazy_(feat._Haillie_Jade): -79755.683138
- 7: 09_-_Drips_(feat._Obie_Trice): -79800.468216
- 8: 09_-_Mmm_Papi_[Main_Version]: -79891.923333
- 9: 06_-_Don't_Want_You_Back: -79933.860554
- 10: 02_-_White_America: -80188.871534
- 11: 13_-_Giddy_Up: -80210.094828
- 12: 10_-_Without_Me: -80229.60496
- 13: 15_-_Phonography_[Main_Version]: -80428.051516
- 14: 17_-_Say_What_You_Say_(feat._Dr._Dre): -80493.740907
- 15: 03_-_Business: -80604.567461

b) Classic Rock (Steely Dan)

Song being searched for: 05_Steely_Dan_-_Home_At_Last

Songs closest to this song:

- 1: 05_Steely_Dan_-_Home_At_Last: -78377.776749
- 2: Steely_Dan_-_Home_At_Last: -79040.130178
- 3: SteelyDan1: -81844.197219
- 4: 01_Steely_Dan_-_Black_Cow: -81908.941216
- 5: 02_Steely_Dan_-_Aja: -82005.614359
- 6: 07_Steely_Dan_-_Josie: -82102.040561
- 7: 06_Steely_Dan_-_I_Got_The_News: -82716.637673
- 8: SteelyDan2: -82755.512876
- 9: 04_Steely_Dan_-_Peg: -82781.06428
- 10: 03_Steely_Dan_-_Deacon_Blues: -83755.671851
- 11: 07-Eric_Clapton-Layla: -84162.728819
- 12: 06-Eric_Clapton-Nobody_Knows_You_When_You're_Down_&_Out: -84598.616681
- 13: 04_-_Yes_-_Roundabout: -84967.870495
- 14: 07_-_Dr._John_-_Right_Place_Wrong_Time: -84979.619995
- 15: 13-Eric_Clapton-Old_Love: -85097.85801

d) R&B (T-Pain)

Song being searched for: 09_Mix'd_Girl

Songs closest to this song:

- 1: 09_Mix'd_Girl: -75667.434432
- 2: 05_5_0'clock_(feat._Lily_Allen_and_Wiz_Khalifa): -79630.147104
- 3: 12_When_I_Come_Home: -79669.205613
- 4: 14_Turn_All_the_Lights_On_(feat._Ne-Yo): -79801.020262
- 5: 02_-_I_Just_Wanna_Be_With_You: -79856.656869
- 6: 13_Best_Love_Song_(feat._Chris_Brown): -79864.081645
- 7: 15_Center_of_the_Stage_(feat._R_Kelly_and_Bei_Maejor): -79870.480037
- 8: 13_-_Giddy_Up: -79888.555924
- 9: 04_Default_Picture: -80048.809021
- 10: 16_Regular_Girl: -80106.410729
- 11: 06_Sho_Time_(Pleasure_Thang): -80129.588706
- 12: 01_-_Tearin'_Up_My_Heart: -80602.266814
- 13: 02_-_Circus_[Main_Version]: -80836.814016
- 14: 08_-_I_Want_You_Back: -80854.505095
- 15: 08_Look_at_Her_Go_(feat._Chris_Brown): -80899.829419

f) Rap (Eminem)

Song being searched for: 07_-_Soldier

Songs closest to this song:

- 1: 07_-_Soldier: -74971.368866
- 2: 14_-_Haillie's_Song: -77266.52674
- 3: 05_-_Square_Dance: -77381.720516
- 4: 09_-_Drips_(feat._Obie_Trice): -77501.950539
- 5: 13_-_Superman_(feat._Dina_Rae): -77623.34898
- 6: 04_-_Cleanin_Out_My_Closet: -77652.86317
- 7: 17_-_Say_What_You_Say_(feat._Dr._Dre): -77823.99845
- 8: 03_-_Business: -77893.344758
- 9: 08_-_Say_Goodbye_Hollywood: -77901.609986
- 10: 10_-_Without_Me: -78133.000954
- 11: 11_-_In_Pieces: -78158.332246
- 12: 16_-_When_The_Music_Stops_(feat._D-12): -78162.271219
- 13: 05_-_God_Must_Have_Spent_A_Little_More_Time_On_You: -78307.462052
- 14: 02_-_White_America: -78330.330947
- 15: 19_-_My_Dad's_Gone_Crazy_(feat._Haillie_Jade): -78339.535255

h) Alternative (Linkin Park)

Song being searched for: 02_-_Don't_Stay

Songs closest to this song:

- 1: 02_-_Don't_Stay: -66422.870663
- 2: 04_-_Bleed_It_Out: -68579.048654
- 3: 07_-_Faint: -68755.398423
- 4: 08_-_Figure.09: -68792.602796
- 5: 05_-_Hit_The_Floor: -68884.471254
- 6: 04_-_Lying_From_You: -69091.109983
- 7: 01_-_Larger_Than_Life: -69263.897448
- 8: 13_-_Numb: -69367.477098
- 9: 08_-_No_More_Sorrow: -69395.665779
- 10: 18_-_Till_I_Collapse_(feat._Nate_Dogg): -69501.140352
- 11: 05_-_Square_Dance: -69752.965882
- 12: 04_-_It's_Gotta_Be_You: -69880.247539
- 13: 19_-_My_Dad's_Gone_Crazy_(feat._Haillie_Jade): -70029.777173
- 14: 14_Turn_All_the_Lights_On_(feat._Ne-Yo): -70144.337005
- 15: 06_-_You_Got_It: -70146.37201

	(a)	(b)	(c)	(d)	(e)	(f)	(g)	(h)	Average
Reciprocal Rank (1/rank of Highest song in same genre)	1	1	1	1	1	1	1	1	1
Precision at Rank 15 (# songs in same genre / # songs retrieved)	1	.93	.93	.8	1	.93	.73	.73	.88
Discounted Cum. Gain (same artist = 10, same genre = 8, other = 0)	194.7	180.8	182.3	194.7	186.0	152.6	154.2	176.3	176.3
Average Precision At Each Point a Song in Same Genre Appears	1	.93	.92	.75	1	.92	.58	.69	.85

Figure V: Clusters found using hierarchical clustering with median similarity criterion.**Classical Music**

Cluster 8:
 Prelude3
 03 - Prelude_in_B_flat_major,_Op.23_No.2
 14 - Prelude_in_E_major,_Op.32_No.3
 22 - Prelude_in_B_major,_Op.32_No.11
 Prelude1
 21 - Prelude_in_B_minor,_Op.32_No.10
 Vladimir_Ashkenazy_-Prelude_in_C_minor,_Op.23_No.7
 08 - Prelude_in_C_minor,_Op.23_No.7
 06 - Prelude_in_G_minor,_Op.23_No.5
 24 - Prelude_in_D_flat_major,_Op.32_No.13
 12 - Prelude_in_C_major,_Op.32_No.1
 16 - Prelude_in_G_major,_Op.32_No.5
 05 - Prelude_in_D_major,_Op.23_No.4
 13 - Prelude_in_B_flat_minor,_Op.32_No.2
 04 - Prelude_in_D_minor,_Op.23_No.3
 15 - Prelude_in_E_minor,_Op.32_No.4
 01 - Prelude_in_C_sharp_minor,_Op.3_No.2
 Prelude2
 10 - Prelude_in_E_flat_minor,_Op.23_No.9
 09 - Prelude_in_A_flat_major,_Op.23_No.8
 07 - Prelude_in_E_flat_major,_Op.23_No.6
 20 - Prelude_in_A_major,_Op.32_No.9
 23 - Prelude_in_G_sharp_minor,_Op.32_No.12
 17 - Prelude_in_F_minor,_Op.32_No.6
 11 - Prelude_in_G_flat_major,_Op.23_No.10
 02 - Prelude_in_F_sharp_minor,_Op.23_No.1
 Vladimir_Ashkenazy_-Prelude_in_F_major,_Op.32_No.7
 18 - Prelude_in_F_major,_Op.32_No.7
 19 - Prelude_in_A_minor,_Op.32_No.8

R&B

Cluster 12:
 04_Default_Picture
 09_Mix'd_Girl
 12_When_I_Come_Home
 06 - Deep_Purple_-Smoke_On_The_Water
 13_Best_Love_Song_(feat._Chris_Brown)
 10_I_Don't_Give_a_Fuck
 14_Turn_All_the_Lights_On_(feat._Ne-Yo)
 08_Look_at_Her_Go_(feat._Chris_Brown)
 15_Center_of_the_Stage_(feat._R._Kelly_and_Bei_Maejor)
 16_Regular_Girl
 11_Drowning_Again_(feat._One_Chance)
 13 - Superman_(feat._Dina_Rae)
 10 - Without_Me
 02 - I_Just_Wanna_Be_With_You
 07_Rock_Bottom
 05_5_O'clock_(feat._Lily_Allen_and_Wiz_Khalifa)
 06_Sho-Time_(Pleasure_Thang)

Rock

Cluster 10:
 09-Eric_Clapton-Walkin'_Blues
 03-Eric_Clapton-Hey_Hey
 02-Eric_Clapton-Before_You_Accuse_Me
 12-Eric_Clapton-Malted_Milk
 11-Eric_Clapton-San_Francisco_Bay_Blues
 10-Eric_Clapton-Alberta
 07-Eric_Clapton-Layla
 06-Eric_Clapton-Nobody_Knows_You_When_You're_Down_&_Out
 04-Eric_Clapton-Tears_In_Heaven
 05-Eric_Clapton-Lonely_Stranger
 14-Eric_Clapton-Rollin'_&_Tumblin'
 08-Eric_Clapton-Running_On_Faith
 01-Eric_Clapton-Signe
 13-Eric_Clapton-Old_Love
 04_Steely_Dan_-_Peg
 10 - The_Bellamy_Brothers_-_Let_Your_Love_Flow
 05_Steely_Dan_-_Home_At_Last
 Steely_Dan_-_Home_At_Last
 SteelyDan1
 01_Steely_Dan_-_Black_Cow

Cluster 6:
 03_Steely_Dan_-_Deacon_Blues
 02_Steely_Dan_-_Aja
 SteelyDan2
 06_Steely_Dan_-_I_Got_The_News
 07_Steely_Dan_-_Josie

Cluster 8:
 02 - Bad_Company_-Feel_Like_Makin'_Love
 08 - Lynyrd_Skynrd_-Sweet_Home_Alabama
 01 - Chicago_-Saturday_In_The_Park
 09 - America_-Sister_Golden_Hair

Pop

Cluster 2:
 07 - Hands_Held_High
 03 - Show_Me_The_Meaning_Of_Being_Lonely
 07 - Don't_Wanna_Lose_You_Now
 12 - The_Perfect_Fan
 04 - For_The_Girl_Who_Has_Everything
 12 - The_Little_Things_Give_You_Away
 09 - Back_To_Your_Heart
 05 - I_Need_You_Tonight

Cluster 14:
 10 - Spanish_Eyes
 11 - No_One_Else_Comes_Close

Figure VI: Clusters found using algorithm described.

Alternative (Linkin' Park)

Cluster 6:
 09 - Valentine's Day
 05 - Hit The Floor
 04 - Bleed It Out
 08 - Figure.09
 12 - Session
 04 - Lying From You
 03 - Leave Out All The Rest
 02 - Don't Stay
 06 - Easier To Run
 06 - What I've Done
 18 - 'Till I Collapse (feat. Nate Dogg)
 03 - Somewhere I Belong
 13 - Numb
 07 - Faint
 05 - Square Dance

Classic Rock

Cluster 4:
 07 - Dr. John - Right Place Wrong Time
 Steely Dan - Home At Last
 08 - Lynyrd Skynrd - Sweet Home Alabama
 05 Steely Dan - Home At Last
 01 - Chicago - Saturday In The Park
 09 - America - Sister Golden Hair
 10 - In Between
 03 Steely Dan - Deacon Blues
 10 - The Bellamy Brothers - Let Your Love Flow
 07 Steely Dan - Josie
 SteelyDan1
 01 Steely Dan - Black Cow
 04 - Yes - Roundabout
 11 - Warren Zevon - Poor Poor Pitiful Me
 04 Steely Dan - Peg

Classical

Cluster 3:
 12 - Prelude in C major, Op.32 No.1
 06 - Prelude in G minor, Op.23 No.5
 05 - Prelude in D major, Op.23 No.4
 01 - Prelude in C sharp minor, Op.3 No.2
 07 - Prelude in E flat major, Op.23 No.6
 21 - Prelude in B minor, Op.32 No.10
 03 - Prelude in B flat major, Op.23 No.2
 02 - Prelude in F sharp minor, Op. 23 No.1
 16 - Prelude in G major, Op.32 No.5
 Vladimir Ashkenazy - Prelude in F major, Op.32 No.7
 18 - Prelude in F major, Op.32 No.7
 22 - Prelude in B major, Op.32 No.11
 09 - Prelude in A flat major, Op.23 No.8
 08 - Prelude in C minor, Op.23 No.7
 Prelude3
 24 - Prelude in D flat major, Op.32 No.13
 14 - Prelude in E major, Op.32 No.3
 10 - Prelude in E flat minor, Op.23 No.9
 19 - Prelude in A minor, Op.32 No.8
 04 - Prelude in D minor, Op.23 No.3
 Prelude1
 17 - Prelude in F minor, Op.32 No.6
 11 - Prelude in G flat major, Op.23 No.10
 Vladimir Ashkenazy - Prelude in C minor, Op.23 No.7
 23 - Prelude in G sharp minor, Op.32 No.12
 13 - Prelude in B flat minor, Op.32 No.2
 Prelude2
 15 - Prelude in E minor, Op.32 No.4
 20 - Prelude in A major, Op.32 No.9

Rock (Clapton)

Cluster 5:
 14 - Eric Clapton - Rollin' & Tumblin'
 04 - Eric Clapton - Tears In Heaven
 11 - Eric Clapton - San Francisco Bay Blues
 10 - Eric Clapton - Alberta
 01 - Eric Clapton - Signe
 05 - Eric Clapton - Lonely Stranger
 05 - God Must Have Spent A Little More Time On You
 08 - Eric Clapton - Running On Faith
 07 - Unusual You [Main Version]
 11 - Drowning Again (feat. One Chance)
 07 - Eric Clapton - Layla
 12 - The Doobie Brothers - Rockin' Down The Highway
 06 - Eric Clapton - Nobody Knows You When You're Down & Out
 13 - Eric Clapton - Old Love
 03 - Eric Clapton - Hey Hey
 09 - Eric Clapton - Walkin' Blues
 02 - Eric Clapton - Before You Accuse Me
 12 - Eric Clapton - Malted Milk

Pop (Boy Bands)

Cluster 1:
 11 - No One Else Comes Close
 07 - Don't Wanna Lose You Now
 09 - Everything I Own
 12 - Sailing
 04 - For The Girl Who Has Everything
 05 - I Need You Tonight
 02 - I Want It That Way
 07 - Hands Held High
 12 - The Perfect Fan
 03 - Out From Under [Main Version]
 09 - Back To Your Heart
 05 - Shadow Of The Day
 04 - Cleanin Out My Closet
 12 - The Little Things Give You Away
 12 - My Baby [Main Version]
 03 - Show Me The Meaning Of Being Lonely
 10 - Spanish Eyes

Pop (Britney Spears)

Cluster 8:
 09 - Drips (feat. Obie Trice)
 15 - Phonography [Main Version]
 14 - Rock Me In [Main Version]
 17 - Say What You Say (feat. Dr. Dre)
 08 - Blur [Main Version]
 02 - Circus [Main Version]
 04 - Kill The Lights [Main Version]
 05 - Shattered Glass [Main Version]
 09 - Mmm Papi [Main Version]
 10 - Mannequin [Main Version]
 03 - Business
 03 - Foreigner - Feels Like The First Time
 02 - White America
 06 - If U Seek Amy [Main Version]
 11 - Lace And Leather [Main Version]
 13 - Radar [Main Version]

R&B (T-Pain)

Cluster 0:
 09 Mix'd Girl
 12 When I Come Home
 15 Center of the Stage (feat. R. Kelly and Bei Maejor)
 03 It's Not You (It's Me) (T-Pain vs. Chuckie) (feat. Pitbull)
 07 Rock Bottom
 05 5 O'Clock (feat. Lily Allen and Wiz Khalifa)
 10 I Don't Give a Fuck
 16 Regular Girl
 05 - Argent - Hold Your Head Up
 17 Nothin' (feat. E-40 and Detail)
 13 Best Love Song (feat. Chris Brown)
 06 Sho-Time (Pleasure Thang)
 04 Default Picture
 01 Bang Bang Pow Pow (feat. Lil Wayne)
 08 Look at Her Go (feat. Chris Brown)